

---

# Efficient Scaling of Bayesian Neural Networks

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1 While Bayesian neural networks (BNNs) have gained popularity for their theoretical  
2 guarantees and robustness, they have yet to see a convincing implementation  
3 at scale. This study investigates a variational inference-based neural architecture  
4 called Variational Density Propagation (VDP) that boasts noise robustness, self-  
5 compression and improved explanations over traditional (deterministic) neural  
6 networks. Due to the large computational burden associated with BNNs, however,  
7 these methods has yet to scale efficiently for real-world problems. In this study, we  
8 simplify the VDP architecture reducing its time and space requirements allowing for  
9 efficient scaling to ImageNet level problems. Additionally, we evaluate the inherent  
10 properties of the VDP method in order to validate the simplified method. Across  
11 all datasets and architectures, our method exhibits exceptional self-compression  
12 capabilities, retaining performance even with over 90% of its parameters pruned.  
13 The method also presents improved visual explanations via saliency maps, sug-  
14 gesting superior explanation quality compared to deterministic models. Lastly, we  
15 employ the VDP method to train a vision transformer on ImageNet-1k, something  
16 that was previously impossible due to the inherent computational constraints of the  
17 method.

## 18 1 Introduction

19 The state of machine learning research has seen tremendous growth, with increasingly complex  
20 and large-scale models being developed [1]. This increase in size and complexity, however, has led  
21 to a notable concern regarding overfitting [2]. As models become larger, they may inadvertently  
22 memorize training data instead of learning to generalize, resulting in poor performance on unseen  
23 or novel data samples. A guiding principle against complexity increase has been the The Minimum  
24 Description Length Principle which states that the best model is one which minimizes the distance  
25 between the model and data and the models description of itself [3]. This principle is the foundation  
26 of stochastic modeling and a benefit of using Bayesian inference.

27 Bayesian inference provides a principled approach for dealing with uncertainty by combining prior  
28 knowledge and observed data to update beliefs about model parameters [4]. In the context of machine  
29 learning, Bayesian Neural Networks (BNNs) extend this framework by incorporating Bayesian  
30 principles into the architecture and training of neural networks [5]. BNNs estimate the posterior  
31 distribution over model parameters, allowing them to capture and quantify uncertainties in both  
32 predictions and model structures. By leveraging Bayesian inference, BNNs offer a more robust and  
33 adaptive learning approach, which can potentially alleviate overfitting and improve generalization  
34 performance in complex machine learning tasks [6].

35 BNNs have demonstrated successful implementations across a wide range of tasks, including com-  
36 puter vision [7–9], speech [10], and natural language processing [11]. BNNs have shown particular  
37 promise in critical decision-making tasks [12–14], due to their ability to capture and model uncertainty  
38 in both parameters and predictions [7]. Several classical techniques have been used to approximate

Bayesian inference for neural networks, such as Laplacian approximation [15], Hamiltonian Monte Carlo [16], and Variational Inference (VI) [6]. Specifically, estimating the parameters of the variational distribution that approximates the posterior of the latent variables given an observation is analogous to averaging model parameters via Monte Carlo sampling [17]. This model averaging inherently reduces variance in computed parameters, thus mitigating overfitting [18]. Moreover, the regularization term of the VI algorithm can lead to extreme self-compression of the model’s parameters [19], further decreasing overfitting.

A major drawback of BNNs is their limited scalability in terms of both model and data size [20]. A high-quality implementation of a VI-based training scheme called Variational Density Propagation (VDP) has emerged, offering state-of-the-art performance on various computer vision tasks, reduced epistemic uncertainty, and an estimate of aleatoric uncertainty [8, 9]. However, the primary limitation of the VDP algorithm is the substantial computational load due to the propagation of the full covariance matrix through layers of the neural network. Propagating these large matrices quickly becomes infeasible in terms of space and time for all but the smallest ResNets, significantly limiting the applicability of the VDP method for problems that require training on datasets larger than CIFAR-10 [21].

In this work, we propose a modification to the VDP method, in which we disregard the off-diagonal terms of the variance-covariance matrix and only propagate the diagonal terms. This simplification enables efficient scaling of these models to larger and more complex datasets, with a theoretical memory requirement only twice that of a traditional deep neural network, while retaining all the benefits of BNNs. Our contributions include:

1. A theoretical justification for the novel variance-only VDP implementation (referred to as VDP++).
2. A comparison of runtime and memory requirements between VDP++ and traditional models.
3. Validation of inherent BNN properties, including uncertainty quantification, self-compression, and robust interpretations or explainability.
4. Implementation of Bayesian VDP Vision Transformers (VDP-ViT) and scaling the model to ImageNet-1k.

## 2 Related Work

**Self-compression:** Since their inception in the early 1990s, BNNs have been developed with the goal of minimizing the information content in the weights of a neural network. This approach is based on the principle that smaller model sizes often lead to more generalizable models [6, 15]. However, contemporary BNN methods such as Bayes by Backprop (BBB) [7] and Dropout CNNs [22] appear to lack this property. In the case of BBB, parameter histograms suggest a less efficient use of parameters compared to traditional deterministic neural networks [7]. To the best of our knowledge, the only modern framework that exhibits this self-compression property is Variational Density Propagation (VDP) [8, 9, 19]. The self-compression characteristic of VDP can be attributed to the propagation of variance information through the network layers [19]. By quantifying parameter uncertainty, the model can selectively target less important parameters during training. In the absence of uncertainty quantification, the network’s utilization of its parameter space is less efficient.

**Uncertainty Quantification:** The total uncertainty in a machine learning model can be characterized by two distinct components. The first is epistemic uncertainty, which arises from the parameters of the model. The second is aleatoric uncertainty, which stems from the data or environment. Uncertainty quantification is becoming increasingly important in high-stakes decision-making, as establishing trust between users and models is crucial for the widespread adoption of such models [23]. To date, the implementation of machine learning models in high-stakes domains, such as medicine, has not achieved substantial user acceptance and confidence [24, 25]. This is primarily due to high false alarm rates [26] and suboptimal test characteristics [27]. VI has been demonstrated to alleviate the impact of epistemic uncertainty by effectively averaging predictions during inference [6, 7, 18]. Furthermore, the output variance associated with a single prediction obtained from a BNN can quantify the level of aleatoric uncertainty in that prediction [8, 28].

**Explainability:** A common approach to interpreting predictions in computer vision models involves calculating sensitivity maps. These gradient-based methods are used to determine the contribution




Figure 1: In order to motivate the needs for a simplification of the VDP algorithm, we measured the: (a) average epoch time, (b) maximum GPU memory usage for each model type training on MNIST with a batch size of 512 and (c) test accuracy on the MNIST dataset with a high level of noise added (higher is better). All measurements were taken using a Quadro RTX 8000 GPU.

of each image pixel to the classification output [29]. However, interpretations derived from deep neural networks, particularly sensitivity maps, are often fragile—meaning they are sensitive to small perturbations in the input or model [30–32]. One potential solution to this issue involves sampling the surrounding input space, adding Gaussian noise, and averaging the resulting interpretations [33]. This averaging technique for reducing uncertainty aligns with our hypothesis on BNNs and weight averaging. It has been demonstrated that the explanations generated by the VDP model outperform those from traditional models [34].

### 3 Methods

**Variance-only Variational Density Propagation (VDP++):** Our work presents a novel simplification of the VDP algorithm developed in [8]. This method utilizes VI and assumes Tensor Normal Distributions (TNDs) defined over the parameters of the neural network in order to propagate the first two moments (mean and covariance) of these TNDs through the layers of a neural network [8]. However, in an experimental setting, propagating covariance proves to be computationally expensive in terms of both time and space. As these networks scale, calculating such large matrices becomes increasingly burdensome, rendering the method impractical. In this work, we have streamlined the VDP algorithm by propagating only the diagonal elements of the covariance matrix, i.e., the variance. This simplification, originally discussed by Hinton [6], is akin to the assumptions of the Naive Bayes classifier [35].

For the sake of brevity, we have omitted the original derivation of the VDP algorithm and only provide the modifications needed to compute the variance-only version [8]. To describe our method, we use a 2-layer fully-connected neural network as an example. In a traditional or deterministic neural network, Equation 1 describes the forward pass of the model.

$$\begin{aligned} z &= Wx + b^{[1]}, \\ a &= f(z), \\ \hat{y} &= g(Va + b^{[2]}), \end{aligned} \tag{1}$$

where  $W \in \mathbb{R}^{j \times k}$  is the weight matrix of layer 1,  $x \in \mathbb{R}^{k \times 1}$  is the input vector,  $b^{[1]} \in \mathbb{R}^{j \times 1}$  is the bias vector in layer 1,  $z \in \mathbb{R}^{j \times 1}$  is the result of the linear operation in layer 1,  $f$  is an arbitrary element-wise non-linear function,  $g$  is an arbitrary non-linear activation function that does not operate element-wise,  $a \in \mathbb{R}^{j \times 1}$  is the result after applying the non-linear activation function in layer 1,  $V \in \mathbb{R}^{l \times j}$  is the weight matrix in layer 2,  $b^{[2]} \in \mathbb{R}^{l \times 1}$  is the bias vector in layer 2,  $\hat{y} \in \mathbb{R}^{l \times 1}$  is the predicted output,  $k$  is the dimensionality of the input vector,  $j$  is the number of nodes in layer 1 and  $l$  is the number of classes to predict.

To propagate the first two moments, several assumptions must be made. First, let us consider  $w_m^\top = m^{\text{th}}$  row of  $W$ ,  $m = 1, 2, \dots, j$  and  $z_m = w_m^\top x + b_m^{[1]}$ ,  $m = 1, 2, \dots, j$ . Next, consider the following assumptions: the input vector  $x$  is deterministic,  $a \sim \mathcal{N}(\mu_a, \Sigma_a)$ ,  $w_m \sim \mathcal{N}(\mu_{w_m}, \Sigma_{w_m})$ ,  $m = 1, 2, \dots, j$ ,  $b_m^{[1]} \sim \mathcal{N}(\mu_{b_m^{[1]}}, \sigma_{b_m^{[1]}}^2)$ ,  $m = 1, 2, \dots, j$ ,  $v_n \sim \mathcal{N}(\mu_{v_n}, \Sigma_{v_n})$ ,  $n = 1, 2, \dots, l$ ,  $b_n^{[2]} \sim \mathcal{N}(\mu_{b_n^{[2]}}, \sigma_{b_n^{[2]}}^2)$ ,  $n = 1, 2, \dots, l$  and the weight vectors  $w_m$ ,  $a$ , bias  $b^{[1]}$  and  $b^{[2]}$  are mutually uncorrelated




Figure 2: Global unstructured pruning was performed on each model type tested on: (a) MNIST and (b) CIFAR-10. This process removes parameters from the model from all layers in by their L1 norm. Here we show that the VDP model types have superior self-compressing properties as they can maintain their performance when greater than 90% of the parameters of the model are removed. Five models were trained for each model and dataset and the 95% confidence interval is displayed as the shaded region.

126 with each other for  $m = 1, 2, \dots, j$ . Based on these assumptions, the elements of  $\mu_z$  and  $\sigma_z^2$  are given  
 127 in Equations 2 and 3.

$$\begin{aligned}
 \mu_{z_m} &= \mathbb{E}[w_m^\top x + b_m^{[1]}], \\
 &= \mathbb{E}[w_m^\top x] + \mathbb{E}[b_m^{[1]}], \\
 &= \mu_{w_m}^\top x + \mu_{b_m^{[1]}}.
 \end{aligned} \tag{2}$$

$$\begin{aligned}
 \sigma_{z_m}^2 &= \text{Var}[w_m^\top x + b_m^{[1]}], \\
 &= x_m^\top \Sigma_{w_m}^2 x_m + \sigma_{b_m^{[1]}}^2, \\
 &= x_m^\top [\sigma_{w_m}^2]^\top + \sigma_{b_m^{[1]}}^2.
 \end{aligned} \tag{3}$$


128 Since we have assumed the weight vectors and the elements of the bias vector to be uncorrelated,  
 129  $\Sigma_{w_p w_q} = 0$  and  $\sigma_{b_p b_q} = 0$  for  $p \neq q$ , where  $p, q = 1, 2, \dots, j$ . Hence, the covariance is zero. To  
 130 propagate the first two moments through an arbitrary element-wise non-linear function, we utilize the  
 131 first-order Taylor series approximation shown in Equations 4 and 5.

$$\begin{aligned}
 a &= f(z), \\
 &= f(\mu_z) + f'(\mu_z)(z - \mu_z) + \dots \\
 &\approx f(\mu_z) + f'(\mu_z)(z - \mu_z), \\
 \mathbb{E}[a] &= \mu_a \approx f(\mu_z), \\
 \mu_a &= f(\mu_z).
 \end{aligned} \tag{4}$$


$$\sigma_a^2 = \sigma_z^2 \odot (f'(\mu_z))^2. \tag{5}$$

132 For the second layer of the network, we can no longer consider the incoming vector,  $\mu_a$ , to be  
 133 deterministic. Additionally, we now need to propagate the incoming variance,  $\sigma_a^2$ . Again, we assume  
 134 the off-diagonal elements of the covariance  $\Sigma_a$  to be 0 and choose to only propagate the variance.  
 135 The mean and variance propagated through the second fully-connected layer are given in Equations 6  
 136 and 7.

$$\begin{aligned}
 \mu_{\tilde{y}} &= \mathbb{E}[v_n^\top a + b_n^{[2]}], \\
 &= \mathbb{E}[v_n^\top] \mathbb{E}[a] + \mathbb{E}[b_n^{[2]}], \\
 &= \mu_{v_n}^\top \mu_a + \mu_{b_n^{[2]}}.
 \end{aligned} \tag{6}$$



(a)



(b)

Figure 3: To validate our model’s robustness to noise, a procedure designed by Dera *et al.* [8] was performed. Gaussian noise was added to the test set of: (a) MNIST and (b) CIFAR-10 to generate sets of images with varying Signal to Noise Ratio (SNR). The output variance of the predicted class was extracted for each prediction and was aggregated over the test set. The output variances were normalized by dividing the variance by the mean variance of the lowest SNR. Five models were trained for each model and dataset and the 95% confidence interval is displayed as the shaded line.

$$\begin{aligned}
\sigma_y^2 &= \text{Var}[v_n^\top a + b_n^{[2]}], \\
&= \text{Var}[v_n^\top a] + \text{Var}[b_n^{[2]}], \\
&= \text{Tr}(\Sigma_v^2 \Sigma_a^2) + \mu_v^\top \Sigma_a^2 \mu_v + \mu_a^\top \Sigma_v^2 \mu_a + \Sigma_{b_n^{[2]}}^2, \\
&= \sigma_v^2 [\sigma_a^2]^\top + \mu_v^2 [\sigma_a^2]^\top + \mu_a^2 [\sigma_v^2]^\top + \sigma_{b_n^{[2]}}^2.
\end{aligned} \tag{7}$$

137 Since the diagonals of the covariance matrix are along the rows of the variance matrix, the trace of  
138 the product of two vectors is equivalent to their inner product.

139 Next, we use a non-linear activation function such as softmax to obtain the predictions of our  
140 model. Since  $g$  does not operate element-wise on our mean and variance, we use a slightly different  
141 Taylor-series to obtain Equations 8 and 9 [36].

$$\mu_{\tilde{y}} \approx g(\mu_{\tilde{y}}), \tag{8}$$

$$\sigma_{\tilde{y}}^2 \approx J_g^2 \sigma_{\tilde{y}}^2, \tag{9}$$

142 where  $J_g$  is the Jacobian matrix of the softmax function  $g$  with respect to  $\tilde{y}$  and calculated at  $\mu_{\tilde{y}}$ .

143 Finally, we use the Evidence Lower Bound (ELBO) function,  $\mathcal{L}(\phi, D)$ , which consists of two  
144 parts: the expected log-likelihood of the training data given the weights, and a regularization term,  
145  $D = \{x^{(i)}, y^{(i)}\}_{i=1}^N$  given by Equation 10.

$$\mathcal{L}(\phi, D) = \mathbb{E}_{q(\phi)}[\log p(D|\phi)] - \text{KL}[q(\phi)|p(\phi)]. \tag{10}$$

146 In Equation 10.  $\phi$  represents the weights  $W$ ,  $V$ , and biases  $b^{[1]}$ ,  $b^{[2]}$ . The expected log-likelihood is  
147 given in Equation 11 and the KL term is given in Equation 12.




Figure 4: To evaluate the effect if any on the quality of explanations produced by BNNs, two metrics were used: Sensitivity Max [32] and Infidelity [37] (lower is better). Perturbations were added to all instances of the test sets of (a) MNIST and (b) CIFAR-10 and the Sensitivity and Infidelity scores were computed and averaged. Five models were trained for each model and dataset and the 95% confidence interval is displayed as the error bar.

$$\begin{aligned}
E_{q(\phi)}[\log p(D|\phi)] &\approx \frac{1}{M} \sum_{m=1}^M \log p(D|\phi), \\
&\approx -\frac{Nl}{2} \log(2\pi) - \frac{1}{M} \sum_{m=1}^M \left[ \frac{N}{2} \log(|\Sigma_{\hat{y}}|) \right. \\
&\quad \left. + \frac{1}{2} \sum_{i=1}^N (y^{(i)} - \mu_{\hat{y}}^{(m)})^\top (\Sigma_{\hat{y}}^{(m)})^{-1} (y^{(i)} - \mu_{\hat{y}}^{(m)}) \right], \\
&\approx -\frac{Nl}{2} \log(2\pi) - \frac{1}{M} \sum_{m=1}^M \left[ \sum_{k=1}^l \log \sigma_{\hat{y}_k}^2 \right. \\
&\quad \left. + \frac{1}{2} \sum_{i=1}^N \left( y^{(i)} - \mu_{\hat{y}}^{(m)} \right)^2 \text{diag}(\sigma_{\hat{y}}^2)^{-1} \right].
\end{aligned} \tag{11}$$

148 In Equation 11,  $y^{(i)}$  is the true label of the  $i^{\text{th}}$  data point,  $N$  is the number of data points and  $M$  is the  
149 number of Monte Carlo samples needed to approximate the expectation by summation.

150 Recall that we have assumed the weight vectors and the elements of the bias vector to be uncorrelated,  
151  $\Sigma_{w_p w_q} = 0$  and  $\Sigma_{b_p b_q} = 0$  for  $p \neq q$ , where  $p, q = 1, 2, \dots, j$ . It follows that  $\Sigma_{\hat{y}}$  is a diagonal matrix  
152 with elements  $\sigma_{\hat{y}_k}^2$ . The log determinant of a diagonal matrix simplifies to the product of the elements  
153 that can be implemented as a sum of the log of the elements,  $\frac{1}{N} \sum_{i=1}^N \left( \sum_{k=1}^l \log \sigma_{\hat{y}_k}^2 \right)$  to prevent  
154 numerical overflow. Additionally, the inverse of a diagonal matrix is equivalent to the reciprocal of  
155 each element,  $\text{diag}(\sigma_{\hat{y}}^2)^{-1}$ .

$$\text{KL}[q(\phi)|p(\phi)] = -\frac{1}{2} \sum_{n=1}^l (j \log \sigma_{v_n}^2 - \|\mu_{v_n}\|_F^2 - j \sigma_{v_n}^2). \tag{12}$$

156 **VDP++ for Convolutional Kernels:** In practice, convolution operations are implemented as matrix  
157 multiplications. Therefore, no additional derivations are needed to propagate the first two moments  
158 through convolutional kernels. For max-pooling, we cannot take the maximum variance. Instead,  
159 we utilize the co-pooling operation. The co-pooling operation is the same for the first moment: the  
160 maximum of the means in the kernel is passed forward. For the variance, we keep only the elements  
161 of the variance that correspond to the maximum means [8].

162 **VDP++ for Residual Connections:** Residual connections are a module used in deep architectures  
163 to resolve the issue of vanishing gradients. These connections propagate parameters from previous

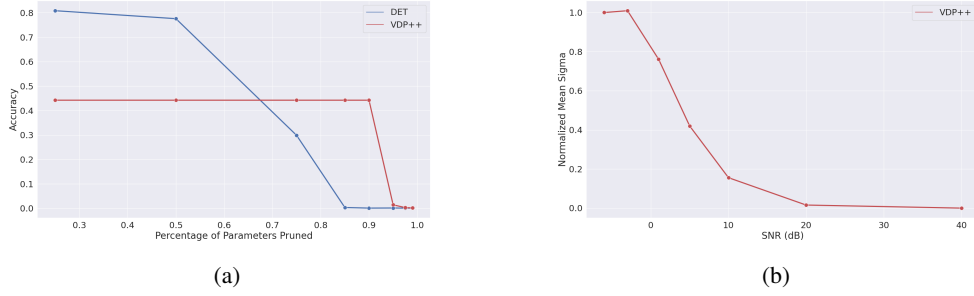


Figure 5: Replication of the pruning experiments (a), and robustness experiments (b) on ImageNet-1k using a VDP-ViT-S/16 architecture. Due to the large hyperparameter selection space and compute time, we were unable to achieve a top-1 accuracy greater than 45%. Although the performance of these models are not high, the VDP retains its self-compressing and noise robustness properties.

164 layers in the neural network by concatenating the output of one layer to the input of the next layer  
 165 [38]. The residual function is effectively a non-elementwise, non-linear function. Therefore, we can  
 166 use the same formulation for Softmax as in Equation 9 [9].

$$\begin{aligned}\mu_{\mathbf{x}_{l+1}} &\approx \mu_{\mathbf{x}_l} + \mathcal{F}(\mu_{\mathbf{x}_l}), \\ \sigma_{\mathbf{x}_{l+1}}^2 &\approx J^2 \sigma_{\mathbf{x}_l}^2.\end{aligned}\tag{13}$$

167 In Equation 13,  $J$  is the Jacobian of  $x_{l+1}$  with respect to  $x_l$  and  $\mathcal{F}$  is the residual function [9].

168 **VDP for Vision Transformers:** A novel contribution in this work, we extend the VDP framework to  
 169 Vision Transformers (ViTs). Given the simple nature of the ViT architecture, to implement models  
 170 like ViT-S/16 and ViT-B/16 [39, 40], we need only implement the layer normalization operation [41]  
 171 for VDP. Similar to the batch normalization formulation for VDP [9], layer normalization for VDP++  
 172 is given in Equation 14 where  $x$  is the input and  $y$  is the output of the layer.

$$\begin{aligned}\mu_y &= \frac{\mu_x - \mu_{LN}}{\sqrt{\text{Var}(\mu_x) + \epsilon}} \odot \gamma + \beta, \\ \sigma_y^2 &= \left( \frac{\mu_{LN}}{\sqrt{\text{Var}(\mu_x) + \epsilon}} \right)^2 \odot \sigma_x^2.\end{aligned}\tag{14}$$

173 **Validation and Robustness:** To validate our approach, we will assess the VDP++ method on the  
 174 MNIST dataset using various back-ends: a deterministic CNN, VDP [8], Bayes by Backprop (BBB)  
 175 [7], and VDP++. On the CIFAR-10 dataset, we will evaluate the deterministic model and VDP++.  
 176 For the MNIST dataset, we employ a LeNet architecture, while for CIFAR-10, we use ResNet-18. In  
 177 order to determine whether any of the backends result in significantly different test accuracy, we will  
 178 repeat training five times and perform a one-way Analysis of Variance (ANOVA) on the results. As  
 179 our approach aims to simplify the original method while preserving its properties, we would like to  
 180 show no significant differences in performance among the methods, i.e.,  $p > \alpha$ .

181 To evaluate the robustness properties of VDP, we will replicate an experiment from the original  
 182 implementation, which entails adding zero-mean Gaussian noise to the test set at varying magnitudes  
 183 and examining the pattern of the normalized output variance [8]. We will use the Spearman rank-  
 184 order correlation measure to compare the two implementations. Additionally, we will compare the  
 185 maximum GPU memory allocation and training time for each formulation at a fixed batch size of 512  
 186 to benchmark each method’s performance.

187 To examine the self-compression properties of VDP++ [19], we will evaluate the performance of  
 188 each backend on the MNIST test set as each model undergoes global pruning by weight magnitude  
 189 (L1 unstructured pruning).

190 Lastly, we will assess the explainability of saliency maps produced by the deterministic model and  
 191 VDP++. The metrics we will utilize to evaluate the quality of the explanations are Infidelity and  
 192 Sensitivity Max [37]. We will employ pair-wise t-tests to determine statistical significance, if any.  
 193 For all statistical analyses, we choose  $\alpha = 0.01$ .

194 **Scaling to Large Datasets:** Remember that the primary limitation of using Bayesian neural networks  
 195 lies in the original formulation’s inability to scale to model sizes larger than ResNet-18. Overcoming  
 196 this drawback would enable the application of BNNs to more complex problems and larger datasets,  
 197 which is essential for advancing the state-of-the-art in machine learning. To that end, we will evaluate  
 198 the performance of VDP++ on the ImageNet-1k dataset using Vision Transformers (ViTs). This  
 199 analysis involves examining the ViT-B/16 [40] and ViT-S/16 [39] architectures. To our knowledge,  
 200 this is the first analysis of BNNs on ImageNet-1k using the ViT architecture.

## 201 4 Results

202 **Time-space Improvements:** We began first by examining the advantages of our method in both  
 203 time and space. On the MNIST dataset, Figure 1 shows a comparison of average epoch time and  
 204 maximum GPU allocation for each method. The original VDP method [8] denoted “VDP Cov” takes  
 205 nearly 10x longer to complete one epoch of training as well as 10x the GPU memory requirement as  
 206 compared to traditional networks (DET) and Bayes-by-Backprop (BBB) [7]. Training on MNIST  
 207 with a batch size of 512 requires nearly 40GB of GPU memory. Our proposed method, VDP++,  
 208 reduces the average epoch time by about 4x over the prior method and reduces the GPU memory  
 209 requirement by 10x.

210 **Self-compression:** The first of the inherent properties of VDP that we investigated was self-  
 211 compression. We began by training 5 models from each method on MNIST (DET, BBB, VDP  
 212 Cov, VDP++) and CIFAR-10 (DET, VDP++). A one-way ANOVA statistic was computed on the test  
 213 statistics of each method and it was found that the performances were not significantly different from  
 214 one another ( $p > 0.01$ ). Next, we iteratively pruned each model using global L1 unstructured pruning  
 215 and computed test statistics. Figure 2 shows the result of this procedure. The VDP Cov and VDP++  
 216 approaches perform similarly on MNIST (Figure 2a). However, VDP++ is able to prune  $> 90\%$  of  
 217 its parameters before the performance begins to drop significantly. This is verified by using a t-test of  
 218 the statistics that compares  $< 90\%$  pruning and  $> 90\%$  pruning. This trend continues to hold for the  
 219 CIFAR-10 experiments (Figure 2b).

220 **Noise Robustness and Uncertainty Quantification:** Figure 1c shows the result of adding zero-mean  
 221 Gaussian noise to the test set of MNIST and examining the test accuracy at the highest level of noise  
 222 (SNR of -6dB). Here the Bayesian models (BBB, VDP Cov, VDP++) all significantly outperform  
 223 the deterministic network (DET). Pair-wise t-tests were performed on deterministic versus each  
 224 Bayesian method at this level of noise and all p-values were significant ( $p < 0.01$ ). Figure 3 shows  
 225 the output variance as a function of SNR normalized by the -6dB value for both MNIST (Figure  
 226 3a) and CIFAR-10 (Figure 3b). When comparing the behavior of VDP-Cov and VDP++ in the  
 227 presence of varying amount of noise (Figure 3a), it was found that there was no significant difference  
 228 (Spearman correlation  $p < 0.01$ ) between the methods.

229 **Explanation Sensitivity:** Due to the high regularization of the network, we postulated that explana-  
 230 tions via saliency maps from Bayesian Networks would be less sensitive to perturbations. Figure 4  
 231 shows the average Sensitivity Max and Infidelity on MNIST (Figure 4a) and CIFAR-10 (Figure 4b)  
 232 (lower is better). Pair-wise t-tests indicate significant differences between methods for all experiments  
 233 ( $p < 0.01$ ). The deterministic method narrowly outperformed VDP++ by Sensitivity Max. However,  
 234 VDP++ significantly outperformed all methods by Infidelity.

235 **ImageNet and VDP-ViT:** To first validate our implementation of VDP-ViT, we trained a small ViT  
 236 on MNIST and confirmed noise robustness using the variance vs SNR experiment. We then began  
 237 parameter sweeps for ViT-B/16 and ViT-S/16 but ultimately decided on only optimizing ViT-S/16  
 238 due to long training times. The training was split across 6 Quadro RTX 8000 GPUs following the  
 239 regime outlined by Beyer *et al.* [39]. Ultimately, we were only able to obtain a model with 45% top-1  
 240 training accuracy given the long training times and large parameter space to search. Figure 5 shows  
 241 the same analyses as above using the underfitted ImageNet-1k model. Even in this state, the VDP++  
 242 model still retains noise robustness and self-compressing properties, remaining consistent with the



243 smaller models. The results for explanation sensitivity were not consistent with our prior results most  
244 likely due to the low performance of the model.

## 245 5 Discussion

246 Our study aimed to investigate the benefits and properties of the proposed novel VDP++ method  
247 compared to the original VDP method and other traditional networks. The results show that VDP++  
248 offers several key advantages, including significant time-space improvements, self-compression,  
249 noise robustness, and uncertainty quantification. The time-space improvements demonstrated by  
250 VDP++ enable more efficient training and reduced GPU memory requirements, which are crucial in  
251 large-scale applications. This addresses the primary limitations of the original VDP method, which  
252 required much longer training times and higher memory consumption.

253 In terms of self-compression and noise robustness, VDP++ performs similarly to the original VDP  
254 method. Additionally, we show that it is able to prune more than 90% of its parameters before  
255 performance significantly declines. This ability to maintain performance despite substantial pruning  
256 suggests that VDP++ learns more efficient representations and could lead to more compact models  
257 that may be beneficial for edge computing.

258 The evaluation of the explanations produced by the various backends revealed unexpected results. Our  
259 original hypothesis was that due to the averaging effect of BNNs, we would have expected to see an  
260 improvement in Sensitivity Max (measures the stability and robustness of the explanation with respect  
261 to small perturbations in the input data) over deterministic models. An improvement in infidelity  
262 would suggest that the explanation approximates the underlying model’s decision-making process  
263 better in BNNs than in deterministic models. The difference in Sensitivity Max, while significant  
264 ( $p < 0.01$ ), was quite small. According to the original work from which it was derived, this score  
265 can be reduced by modifying the formulation of the saliency map, while infidelity cannot [37]. This  
266 yields promising results for the use of BNNs in areas where interpretability and explainability are  
267 desirable.

268 During our testing, we found that the balance between the two terms of the loss function in Equation  
269 10 had the largest effect on the convergence and subsequent properties of our models. Since the  
270 KL term of Equation 12 is largely dominated by matrix norms, in larger models especially, this  
271 term tends to dominate. When this happens, the model compresses itself at the cost of performance.  
272 Conversely, when the KL term is scaled down too much (i.e.  $\mathbb{E}_{q(\phi)}[\log p(D|\phi)] \gg \text{KL}[q(\phi)||p(\phi)]$ ),  
273 the model overfits and the resulting models lack the inherent properties like self-compression and  
274 noise robustness, effectively becoming a deterministic model. To mitigate this behavior, we provided  
275 scaling terms to our hyperparameter optimizer and searched outright. This proved to be the fastest  
276 and most reliable method for training VDP++ models with high performance while retaining their  
277 inherent properties.

278 For classification problems such as ImageNet-1k, the Negative Log-Likelihood (NLL) term in the  
279 ELBO loss function became a bottleneck. In practice, the NLL term uses one-hot encoded labels.  
280 This led to vanishing gradients and low model performance. To alleviate this problem, we changed  
281 the NLL term to categorical cross-entropy for our ImageNet-1k experiments. We confirmed that this  
282 change did not negatively affect our prior results.

## 283 6 Limitations and Future Work

284 Our limitation in this study was a lack of computing resources. Given the increased parameter  
285 space we needed to search for our ImageNet-1k experiment, it became infeasible to expect high  
286 model performance in a short period of time. It is also worth noting that the original ViT formula-  
287 tion made use of the JFT-300M dataset for pretraining, which is not publicly available. We have  
288 substantially reduced the computational burden of the original VDP method while preserving its  
289 desirable properties. In tandem with our preliminary results on ImageNet-1k, scaling this approach  
290 to larger datasets is both possible and feasible given adequate computing resources and time. The  
291 robustness and improved explanatory capabilities of the VDP framework make it highly attractive  
292 for high-risk domains such as healthcare. Furthermore, the self-compressing properties render this  
293 technique suitable for applications with varying resource constraints, such as edge computing and  
294 large language models where computational resources are highly valuable.

## References

- [1] OpenAI, “Gpt-4 technical report,” 2023.
- [2] N. Carlini, F. Tramer, E. Wallace, M. Jagielski, A. Herbert-Voss, K. Lee, A. Roberts, T. B. Brown, D. Song, U. Erlingsson, *et al.*, “Extracting training data from large language models,” in *USENIX Security Symposium*, vol. 6, 2021.
- [3] J. Rissanen, “Stochastic complexity and modeling,” *The annals of statistics*, pp. 1080–1100, 1986.
- [4] G. E. Box and G. C. Tiao, *Bayesian inference in statistical analysis*. John Wiley & Sons, 2011.
- [5] D. J. MacKay, “The evidence framework applied to classification networks,” *Neural computation*, vol. 4, no. 5, pp. 720–736, 1992.
- [6] G. E. Hinton and D. Van Camp, “Keeping the neural networks simple by minimizing the description length of the weights,” in *Proceedings of the sixth annual conference on Computational learning theory*, pp. 5–13, 1993.
- [7] C. Blundell, J. Cornebise, K. Kavukcuoglu, and D. Wierstra, “Weight uncertainty in neural network,” in *International conference on machine learning*, pp. 1613–1622, PMLR, 2015.
- [8] D. Dera, G. Rasool, and N. Bouaynaya, “Extended variational inference for propagating uncertainty in convolutional neural networks,” in *2019 IEEE 29th International Workshop on Machine Learning for Signal Processing (MLSP)*, pp. 1–6, IEEE, 2019.
- [9] D. Dera, N. C. Bouaynaya, G. Rasool, R. Shterenberg, and H. M. Fathallah-Shaykh, “Premium-cnn: Propagating uncertainty towards robust convolutional neural networks,” *IEEE Transactions on Signal Processing*, vol. 69, pp. 4669–4684, 2021.
- [10] X. Xie, X. Liu, T. Lee, S. Hu, and L. Wang, “Blhuc: Bayesian learning of hidden unit contributions for deep neural network speaker adaptation,” in *ICASSP 2019-2019 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pp. 5711–5715, IEEE, 2019.
- [11] Y. Xiao and W. Y. Wang, “Quantifying uncertainties in natural language processing tasks,” in *Proceedings of the AAAI conference on artificial intelligence*, vol. 33, pp. 7322–7329, 2019.
- [12] A. Bate, M. Lindquist, I. R. Edwards, S. Olsson, R. Orre, A. Lansner, and R. M. De Freitas, “A bayesian neural network method for adverse drug reaction signal generation,” *European journal of clinical pharmacology*, vol. 54, no. 4, pp. 315–321, 1998.
- [13] Y. Kwon, J.-H. Won, B. J. Kim, and M. C. Paik, “Uncertainty quantification using bayesian neural networks in classification: Application to ischemic stroke lesion segmentation,” in *Medical Imaging with Deep Learning*, 2018.
- [14] Y. Kwon, J.-H. Won, B. J. Kim, and M. C. Paik, “Uncertainty quantification using bayesian neural networks in classification: Application to biomedical image segmentation,” *Computational Statistics & Data Analysis*, vol. 142, p. 106816, 2020.
- [15] D. J. MacKay, “A practical bayesian framework for backpropagation networks,” *Neural computation*, vol. 4, no. 3, pp. 448–472, 1992.
- [16] R. M. Neal, *Bayesian learning for neural networks*, vol. 118. Springer Science & Business Media, 2012.
- [17] D. M. Blei, A. Kucukelbir, and J. D. McAuliffe, “Variational inference: A review for statisticians,” *Journal of the American statistical Association*, vol. 112, no. 518, pp. 859–877, 2017.
- [18] E. Hüllermeier and W. Waegeman, “Aleatoric and epistemic uncertainty in machine learning: An introduction to concepts and methods,” *Machine Learning*, vol. 110, no. 3, pp. 457–506, 2021.
- [19] G. Carannante, D. Dera, G. Rasool, and N. C. Bouaynaya, “Self-compression in bayesian neural networks,” in *2020 IEEE 30th International Workshop on Machine Learning for Signal Processing (MLSP)*, pp. 1–6, IEEE, 2020.

- [20] J. M. Hernández-Lobato and R. Adams, “Probabilistic backpropagation for scalable learning of bayesian neural networks,” in *International conference on machine learning*, pp. 1861–1869, PMLR, 2015.
- [21] A. Krizhevsky, G. Hinton, *et al.*, “Learning multiple layers of features from tiny images,” 2009.
- [22] Y. Gal and Z. Ghahramani, “Bayesian convolutional neural networks with bernoulli approximate variational inference,” *arXiv preprint arXiv:1506.02158*, 2015.
- [23] E. Begoli, T. Bhattacharya, and D. Kusnezov, “The need for uncertainty quantification in machine-assisted medical decision making,” *Nature Machine Intelligence*, vol. 1, no. 1, pp. 20–23, 2019.
- [24] A. D. Bedoya, M. E. Clement, M. Phelan, R. C. Steorts, C. O’Brien, and B. A. Goldstein, “Minimal impact of implemented early warning score and best practice alert for patient deterioration,” *Critical care medicine*, vol. 47, no. 1, p. 49, 2019.
- [25] J. L. Guidi, K. Clark, M. T. Upton, H. Faust, C. A. Umscheid, M. B. Lane-Fall, M. E. Mikkelsen, W. D. Schweickert, C. A. Vanzandbergen, J. Betesh, *et al.*, “Clinician perception of the effectiveness of an automated early warning and response system for sepsis in an academic medical center,” *Annals of the American Thoracic Society*, vol. 12, no. 10, pp. 1514–1519, 2015.
- [26] J. C. Ginestra, H. M. Giannini, W. D. Schweickert, L. Meadows, M. J. Lynch, K. Pavan, C. J. Chivers, M. Draugelis, P. J. Donnelly, B. D. Fuchs, *et al.*, “Clinician perception of a machine learning-based early warning system designed to predict severe sepsis and septic shock,” *Critical care medicine*, vol. 47, no. 11, pp. 1477–1484, 2019.
- [27] M. Dewan, N. Muthu, E. Shelov, C. P. Bonafide, P. Brady, D. Davis, E. S. Kirkendall, D. Niles, R. M. Sutton, D. Traynor, *et al.*, “Performance of a clinical decision support tool to identify picu patients at high risk for clinical deterioration,” *Pediatric Critical Care Medicine*, 2019.
- [28] A. Y. Foong, Y. Li, J. M. Hernández-Lobato, and R. E. Turner, “‘in-between’ uncertainty in bayesian neural networks,” *arXiv preprint arXiv:1906.11537*, 2019.
- [29] K. Simonyan, A. Vedaldi, and A. Zisserman, “Deep inside convolutional networks: Visualising image classification models and saliency maps,” *arXiv preprint arXiv:1312.6034*, 2013.
- [30] D. Alvarez-Melis and T. S. Jaakkola, “On the robustness of interpretability methods,” *arXiv preprint arXiv:1806.08049*, 2018.
- [31] J. Adebayo, J. Gilmer, M. Muelly, I. Goodfellow, M. Hardt, and B. Kim, “Sanity checks for saliency maps,” *Advances in neural information processing systems*, vol. 31, 2018.
- [32] A. Ghorbani, A. Abid, and J. Zou, “Interpretation of neural networks is fragile,” in *Proceedings of the AAAI conference on artificial intelligence*, vol. 33, pp. 3681–3688, 2019.
- [33] D. Smilkov, N. Thorat, B. Kim, F. Viégas, and M. Wattenberg, “Smoothgrad: removing noise by adding noise,” *arXiv preprint arXiv:1706.03825*, 2017.
- [34] I. E. Nielsen, R. P. Ramachandran, N. Bouaynaya, H. M. Fathallah-Shaykh, and G. Rasool, “Evalattai: A holistic approach to evaluating attribution maps in robust and non-robust models,” *arXiv preprint arXiv:2303.08866*, 2023.
- [35] I. Rish *et al.*, “An empirical study of the naive bayes classifier,” in *IJCAI 2001 workshop on empirical methods in artificial intelligence*, vol. 3, pp. 41–46, 2001.
- [36] D. Simon, *Optimal state estimation: Kalman, H infinity, and nonlinear approaches*. John Wiley & Sons, 2006.
- [37] C.-K. Yeh, C.-Y. Hsieh, A. Suggala, D. I. Inouye, and P. K. Ravikumar, “On the (in) fidelity and sensitivity of explanations,” *Advances in Neural Information Processing Systems*, vol. 32, 2019.
- [38] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 770–778, 2016.

- 389 [39] L. Beyer, X. Zhai, and A. Kolesnikov, “Better plain vit baselines for imagenet-1k,” *arXiv*  
390 *preprint arXiv:2205.01580*, 2022.
- 391 [40] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani,  
392 M. Minderer, G. Heigold, S. Gelly, *et al.*, “An image is worth 16x16 words: Transformers for  
393 image recognition at scale,” *arXiv preprint arXiv:2010.11929*, 2020.
- 394 [41] J. L. Ba, J. R. Kiros, and G. E. Hinton, “Layer normalization,” *arXiv preprint arXiv:1607.06450*,  
395 2016.